

Requirements

We used python 3.9.12 for development, on Linux Ubuntu environment. The sources contain the following files:

`requirements.txt` – the packages specification

`requirements.conda.txt` – contains full list of packages extracted by “conda list”

`requirements.pip.txt` – contains full list of packages extracted by “pip freeze”

We ran the code on Ubuntu 20.04.2 LTS (GNU/Linux 5.4.0-124-generic x86_64), with AMD EPYC 7502P Processor

Directory Structure

Our code is based on the challenge’s GitHub repository. We made a few changes to the original files in order to speed up the data loading process by providing pickle data generator. We also made a small modification to the RouteNet model so that it would have an option of returning internal features.

`common/` - some helper code

`config/` - YAML configuration files used for generating data samples

`data_exploration/` - code utilities for analyzing data samples

`datagen/` - code for generating data samples

`generated_datasets/` - placeholder for generated samples

`generated_datasets_pkl/` - placeholder for pickle versions of generated samples

`validation_dataset/` - placeholder for validation data

`validation_dataset_pkl/` - placeholder for pickle version of validation data

`oracle_models/` - placeholder for oracle models

`notebooks/` - Jupyter notebooks for data and results exploration

`random_train/` - scripts for generating and testing training datasets of 100 samples

`RouteNet_Fermi/` - code from original repo (with some modifications)

`RouteNet_utils/` - additional utilities

Besides the above, the top level contains data extraction and training scripts.

Additional Resources

In addition to the code, we provide for [download here](#) a tar file with the following artifacts:

`generated_datasets/` - all the samples we generated (~270K) with both the input to simulator and its outputs

`generated_datasets_pkl/` - all the generated samples in pickle format, each holding RouteNet input structure

`validation_dataset/` - the validation data as provided in the challenge

`validation_dataset_pkl/` - the validation data converted to pickle format, each holding RouteNet input structure

`oracle_models/` - the two oracle models we used to cluster the training samples. Each model’s subdirectory contains the training logs, the best checkpoint, embedding files for all samples and sample losses measured by the model on all the samples.

best_submission/ - contains the submission which got us the best result, including its original setup and pickle files list

Notes:

- If using our provided models and data, place them each at its appropriate location in this tree. For example, place the contents of *generated_samples_pkl/* from our package into the directory *generated_samples_pkl/* in the source tree.
- If not using our provided models and data, and generating everything from scratch, keep these placeholder directories empty. The given instructions should generate their outputs at these locations.
- Throughout this document we refer to the root of directory structure as **\$src**. If using Linux bash it can be set via “*export src=/path/to/sources/root*”. Also, in the following bash scripts please update this variable to point at the correct location.
- If using our provided models and data, you can skip most of this document and go directly to “Clustering and Sampling” section.
- In addition, add **\$src** to your **PYTHONPATH** environment variable, to prevent import problems (e.g. “*export PYTHONPATH=\$src:\$PYTHONPATH*” if using Linux bash).

Generating the datasets

We generated our dataset incrementally, trying out different configurations. The data generation parameters and different configurations are located under *config/datagen* in the yaml files. These files specify topology creation, routing, link bandwidth, buffer sizes, types of queues, packet size, timing distributions, etc. For generating data, we use [Hydra](#) configuration management library to mix and match different configuration options at the command line. Our data is broken into twenty-eight sub-datasets, each generated from a particular distribution. It is worth noting that each generated data folder contains a *.hydra/config.yaml* file. This file lists all the parameters used to generate samples in that specific folder.

1. Generating input data for OMNet

All the datasets can be generated by running ***datagen/generate_all.sh***

This is a bash script that runs our data generation tool several times and creates multiple dataset directories. The script invokes python, so it expects the right python environment to be set.

Before running the script, edit it to replace variable “src” at the beginning of the script to point to the base of the code tree.

```
% chmod +x $src/datagen/generate_all.sh
% $src/datagen/generate_all.sh
```

After the script is finished, you will find the {0/ .. 15/, hard1/ .. hard6/, hard5_small/ .. hard5_small6/} (a total of twenty-eight datasets) folders under *generated_datasets/* directory.

2. Running OMNet

After *generate_all.sh* completes, OMNet needs to be run on all the directories to generate simulation results. This step is described in detail in the challenge description.

3. Converting to pickle format

After OMNet simulation completes, we run *RouteNet_utils/data_utils.py* to create pickle files from the generated data that contain RouteNet compatible input of all samples. This speeds up the data loading process during the validation and training steps. The input to this script should be the top level directory containing a dataset. The script recursively finds all the generated samples and creates a similar directory tree in the target directory which contains a separate pickle file for each sample.

```
python $src/RouteNet_utils/data_utils.py --input_data_dir
$src/generated_datasets --output_data_dir $src/generated_datasets_pkl
```

Same process needs to be repeated also for the validation dataset:

```
python $src/RouteNet_utils/data_utils.py --input_data_dir
$src/validation_dataset --output_data_dir $src/validation_dataset_pkl
```

After these steps, the directory structure in *generated_datasets_pkl* and *validation_dataset_pkl* should resemble the structure in the original data directories. With a pickle file for every sample, replacing the original tar.gz.

Creating Oracle Models

We provide two trained oracle models that were used for the final submission. We also provide the sample losses and embeddings extracted using these oracles. Below are the steps to reproduce the oracles, losses and embeddings.

1. Training the Oracles

The oracles can be trained using the provided scripts *train_oracle_589.sh* and *train_oracle_615.sh*. Since the oracle is trained only on a subset of the entire dataset, each of these scripts will copy the relevant data subset into the training directory (see the script for exact subset specification).

Before running the scripts, edit and modify the variable *\$src* inside the script to point to the code root directory.

```
% chmod +x $src/train_oracle_589.sh
% $src/train_oracle_589.sh
```

```
% chmod +x $src/train_oracle_615.sh
% $src/train_oracle_615.sh
```

The oracle models are saved to *\$src/oracle_models/*.

2. Extracting sample losses and embeddings

After the oracle is trained we use it to extract losses and embeddings for each of the samples in the validation and training datasets. They will be used in later steps.

To extract these using an oracle checkpoint run *oracle_setup.py* script once for validation set and once for all the generated samples. For 6.15 oracle, run from *\$src/oracle_models/6.15* the following commands (assuming the best checkpoint is *43-6.15.index*):

for validation samples:

```
python $src/oracle_setup.py -ckpt <path/to/modelCheckpoints/43-6.15> -  
name val -data $src/validation_dataset.pkl -o .
```

for generated samples:

```
python $src/oracle_setup.py -ckpt <path/to/modelCheckpoints/43-6.15> -  
name train -data $src/generated_datasets.pkl -o .
```

Note that **the checkpoint name is given without the .index suffix**. Also note that this step may take long time to process all the samples.

After the above two commands, the script should have created the following files under the target (-o) directory:

Embeddings:

sample_embeddings_43-6.15/val_min_max_mean.pkl

sample_embeddings_43-6.15/train_min_max_mean.pkl

csv files containing losses for each sample:

eval/val_sample_loss_43-6.15.csv

eval/train_sample_loss_43-6.15.csv

Run a similar command from *\$src/oracle_models/5.89* with the appropriate checkpoint to generate similar files for the second oracle. (Both are necessary for next steps)

Clustering and Sampling 100-sample subsets

As described in the solution document, given the oracle embeddings and validation losses, our code clusters the generated samples into several groups. Then it selects the top-k samples from each group, after which it chooses 100 samples from the selected sub-groups.

The script *\$src/random_train/clustering_train.py* generates multiple sets of 100 samples from the cluster pool. It can also be used to train each of the sets after sampling. The following creates a subdirectory for each sampled set, with file *sample.txt* that holds the list of the chosen pickle samples:

```
python $src/random_train/clustering_train.py --n_trainings <#sampled sets> -d  
$src/generated_datasets.pkl --save_path <path_to_output_dir> --embed_dir
```

```
$src/oracle_models/6.15/sample_embeddings_XXXXX  
$src/oracle_models/5.89/sample_embeddings_XXXXX
```

Replace above *sample_embeddings_XXXXX* with the correct directory name containing the embeddings. The *save_path* parameter should include the full path, e.g., *\$src/random_sets*.

Note - If using the data, oracles and embedding files that we provided, the clustering algorithm generates from the 270K samples a small sampling pool of 401 samples divided into 88 groups. For reference, this pool is saved in the output directory as file “topk_samples_pool.pkl”. User can verify that the 100 samples from our best submission are contained in this pool.

In order to also train on each set and evaluate the results, the following command can be used:

```
python $src/random_train/clustering_train.py --train True --n_trainings  
<#sampled sets> -d $src/generated_datasets.pkl -t $src/validation_dataset.pkl  
--save_path <path_to_output_dir> --embed_dir  
$src/oracle_models/6.15/sample_embeddings_XXXXX  
$src/oracle_models/5.89/sample_embeddings_XXXXX -w_train <#train_cpu_workers>  
-w_val <#val_cpu_workers> --ngpus_val <#gpus> --ckpt_weights_dir  
$src/RouteNet_Fermi/initial_weights/initial_weights
```

In this case, the script will launch a number of parallel workers as specified – each for a different experiment. If *ngpus_val* is specified and its value is greater than 0, the script will run validation on the available GPUs.

In addition, the following flags can be used to tweak the resulting sample generation:

--topk - flag sets the number of top-k nearest neighbors used for sampling.

--permute - specifies whether to shuffle the set after sampling.

In our experiments we got the best models setting these hyperparameters to **--topk 5 --permute True** or **--topk 3 --permute False**.

In order to collect the results from all the experiments, the following python snippet can be used:

```
from random_train.utils import get_eval_loss_multi  
res_df = get_eval_loss_multi(<results_dir>)
```

It returns a Pandas data-frame, where each row describes a single experiment with the following information: loss per epoch, final loss, best loss, best epoch, data directory name and full path of data directory.

Training using tar.gz instead of pickle

During the challenge we trained and validated our solutions using our code, which as shown above receives its input as a text file with 100 pickle paths, rather than a tar.gz directory.

The following script converts a list of pickle files into a proper directory dataset with the appropriate 100 tar.gz files. We used it mainly before submitting to evaluation server, but it can also be used to create a valid 100-sample data directory for training the provided GNN-challenge model. The script copies the

tar.gz versions of the 100 pkl samples into a single folder and reorders them to match datanetAPI data loader.

```
python $src/prepare_submission_data.py -p <path_to_samples_file> -n  
<experiment_name>
```

The resulting dataset will be placed in *\$src/data_tars* folder.

When making our submissions, before calling `generate_submission.py` we run the above script on the samples file which gave us best result and it creates a data directory containing all the relevant files (gz, routing, topologies) for those 100 samples.

One last note – during training datanetAPI code always returns samples in the same order, given by *“random.Random(1234).shuffle(tuple_files)”* (line 798 in *datanetAPI.py*). In contrast our code always returns samples in same order, but it is the order of files given in our text file. Therefore, when converting from our code to original training, `prepare_submission_data.py` also reorders the files so that in original code, the files will be served to training in exactly same order as in our code.